

A Study on “Vision Transformers Need Registers”

Juntang Wang¹

¹Duke Kunshan University (DKU), China

April 25, 2025

“Vision Transformers Need Registers”: Overview

Darcet et al. 2024, *ICLR*, Vision Transformers Need Registers

- **Paper Link:** <https://openreview.net/forum?id=2dnO3LLiJ1>
- **Repo. Link:** <https://github.com/kyegomez/Vit-RGTS> (for replication)
- **TL;DR:**
 - Vision Transformers (ViTs) exhibit **artifacts**:

high-norm tokens in low-information regions
 - In both supervised and self-supervised; disrupt downstream tasks.
 - **Registers**: additional tokens; only for internal computation.
 - This approach:
 - Achieves **SOTA** on dense prediction tasks.
 - Enables **object discovery** in large models.
 - Produces smoother **feature and attention maps**.

1. Problems

- **Motivation**
- **Analysis**

2. Method

3. Experiments

4. Discussion

Motivation

Big picture:

Embedding images into generic features that can serve multiple purposes in CV

Earlier, e.g. SIFT by Lowe 2004, *IJCV*.

Today, “Labels are expensive”: Pretrain → extract features.

Motivation (cont.)

To bring us on the same page:




Figure 1: Self-attention from a ViT trained with no supervision (Caron et al. 2021).

- Object segmentations from the self-attention¹.

¹Self-attention of the [CLS] query produced by the last attention layer (**attention map**).

Motivation (cont.)

The paper's motivation:

Observed: In LOST, DINOv2 *worse* than DINO. Analyze, fix these ViTs.

LOST²: DINO features for Obj. detection³. → A clean attention map? .

²Siméoni et al. 2021, *BMVC*; Seed points → full object regions.

³Not using attention maps (DINO-seg), but patch features.

Motivation (cont.)




Figure 2: *All models but DINO* exhibit peaky outlier values in the attention maps (Darcet et al. 2024).

Artifacts are high-norm outlier tokens






Figure 3: Comparison of local feature norms (Darcet et al. 2024)⁴

⁴150 is the threshold used for artifacts in the paper if not otherwise specified.


Outliers appear during the training of large models



(a) Norms along layers.



(b) Norms along iterations.




(c) Norms across model size.

Figure 4: Illustration of several properties of outlier tokens (Darcet et al. 2024)⁵.

⁵40-layer DINOv2 ViT-g model.

Analysis (cont.)

High-norm tokens appear where patch information is redundant (Fig. 5a)



(a) Cosine similarity to neighbors.

	position prediction		reconstruction
	top-1 acc	avg. distance ↓	L2 error ↓
normal	41.7	0.79	18.38
outlier	22.8	5.09	25.23

(b) Linear probing for local information.

Figure 5: (a): Distribution of cosine similarity between input patches and their 4 neighbors. (b): Local information probing on normal and outlier patch tokens (Darcet et al. 2024).

Analysis (cont.)

High-norm tokens hold little local information (Tab. 5b)

Artifacts hold global information (Tab. 1)

Table 1: Image classification via linear probing on normal and outlier patch tokens (Darcet et al. 2024)⁶.

	IN1k	P205	Airc.	CF10	CF100	CUB	Cal101	Cars	DTD	Flow.	Food	Pets	SUN	VOC
[[CLS]]	86.0	66.4	87.3	99.4	94.5	91.3	<u>96.9</u>	91.5	85.2	99.7	94.7	96.9	78.6	<u>89.1</u>
normal	65.8	53.1	17.1	97.1	81.3	18.6	73.2	10.8	63.1	59.5	74.2	47.8	37.7	70.8
outlier	<u>69.0</u>	<u>55.1</u>	<u>79.1</u>	<u>99.3</u>	<u>93.7</u>	<u>84.9</u>	97.6	<u>85.2</u>	<u>84.9</u>	<u>99.6</u>	<u>93.5</u>	<u>94.1</u>	<u>78.5</u>	89.7

⁶Also report the accuracy of classifiers learnt on the class token.

1. Problems

2. Method

- Hypothesis
- Remediation

3. Experiments

4. Discussion

Hypothesis

In ViT

- Large, long iterations → artifacts.
- *Artifacts are not inherently bad* ⁷.
- However, LOST suffers ⁸.

⁷They often appear in low-information regions and their local role can be recovered by surrounding tokens.

⁸And other tasks that rely on *clean local features*.

Hypothesis (cont.)

Inspired by Bulatov et al. 2022, add *registers* tokens⁹:

Additional learnable input tokens, discarded in output.

This should isolate the global context, preserving clean, local patch features.

⁹First proposed for remediate translation in NLP.

Remediation




Figure 6: Illustration of the proposed remediation and resulting model. (Darcet et al. 2024).

1. Problems

2. Method

3. Experiments

- **Evaluation**
- **Ablation study**
- **Other experiments**

4. Discussion

Experimental Setup

Table 2: Summary of models used in this paper’s experiments.

Method	Sup.	Dataset	Arch.	Paper
DEIT-III ¹⁰	Label	ImageNet-22k	ViT-B	Touvron et al. 2022
OpenCLIP ¹¹	Text	Shutterstock	ViT-B/16	Ilharco et al. 2021
DINOv2 ¹²	Self	ImageNet-22k	ViT-L	Oquab et al. 2024

¹⁰<https://github.com/facebookresearch/deit>

¹¹https://github.com/mlfoundations/open_clip

¹²<https://github.com/facebookresearch/dinov2>

Register tokens effectively remove the norm outliers that were present previously.




Figure 7: Effect of register tokens on the distribution of output norms (Darcet et al. 2024).

Using registers not only *does not degrade performance*, but even:

improves performance by a slight margin in some cases (Tab. 3)

Evaluation (cont.)

Table 3: Evaluation of downstream performance of the models that we trained, with and without registers.

	ImageNet	ADE20k	NYUd
	Top-1	mIoU	rmse ↓
DeiT-III	84.7	38.9	0.511
DeiT-III+reg	84.7	39.1	0.512
OpenCLIP	78.2	26.6	0.702
OpenCLIP+reg	78.1	26.7	0.661
DINOv2	84.3	46.6	0.378
DINOv2+reg	84.8	47.9	0.366

(a) Linear evaluation with frozen features.

	ImageNet
	Top-1
OpenCLIP	59.9
OpenCLIP+reg	60.1

(b) Zero-shot classification.

To remove artifacts¹³::

One register is sufficient, more leads to improved downstream performance. (Fig. 8)

¹³Evaluated on three tasks (ImageNet, ADE-20k and NYUd)

Ablation study (cont.)




Figure 8: Ablation of the the number of register tokens used with a DINOv2 model (Darcet et al. 2024).

Other experiments - Object discovery

Registers *significantly* improve object discovery performance (Tab. 4)

	VOC 2007	VOC 2012	COCO 20k
DeiT-III	11.7	13.1	10.7
DeiT-III+reg	27.1	32.7	25.1
OpenCLIP	38.8	44.3	31.0
OpenCLIP+reg	37.1	42.0	27.9
DINOv2	35.3	40.2	26.9
DINOv2+reg	55.4	60.0	42.0

Table 4: Unsupervised Object Discovery using LOST on models with and without registers (Darcet et al. 2024)¹⁴.

¹⁴Performance measured using corloc.

Other experiments - Qualitative evaluation

Similarly to slot attention (Locatello et al. 2020)¹⁵:

Register tokens attend to different parts of the feature map.

(It's a pity that this behaviour is not explored further in the paper.)




Figure 9: Comparison of the attention maps of the [CLS] and register tokens (Darcet et al. 2024).

¹⁵This behaviour was never required from the model, and emerged naturally from training.

1. Problems

2. Method

3. Experiments

4. Discussion

- **Conclusion**
- **Further work**

- *Artifacts.*
- — likely due to iterations and larger model.
- A possible interpretation: the model reuse low-info areas to store *global context*.
- *Artifacts are not inherently bad*, but can hurt LOST¹⁶.
- *Register tokens* act as a decoupling mechanism.
- Different register tokens seem to focus on *different objects*.

¹⁶And other tasks that rely on clean feature maps

Further work

- *One more benchmark.* Object discovery
 - *Open Images, Object365, WIDER FACE, DeepFashion2, KITTI, AIOD*
- *Improve the model / one more benchmark:*

Figure 10: Preliminary results of our study.

Tiny ImageNet ¹⁷	
Top-1	
DeiT-B/16-D	91.0
DeiT-B/16-D+reg	90.8

(a) Linear evaluation with frozen features.


ImageNet	
Top-1	
OC(ResNet50)	59.6
OC(ResNet50)+reg	60.0
OC(ViT-B/32)	63.2
OC(ViT-B/32)+reg	61.4

(a) Zero-Shot Classification.

¹⁷<https://paperswithcode.com/dataset/tiny-imagenet>

Further work (cont.)

- Different registers may attend to different objects — test this more directly.
 - CLEVR, COCO Panoptic, GQA
- *Move beyond ViT.* e.g. Sequence data, matrix data¹⁸
- *Move beyond ViT.* e.g. Configurations from Multi-Resolution Graph-based Clustering¹⁹



¹⁸<https://www.kaggle.com/competitions/stanford-rna-3d-folding>

¹⁹<https://www.qqgjyx.com/mheatmap/>

Acknowledgements

We thank Prof. Dongmian Zou
for his guidance and support.

This work is supported in part by
◇ STATS 403: Deep Learning^a

^aDKU UG Course, Instructor: Prof. Dongmian Zou

References I

- Bulatov, Aydar, Yuri Kuratov, and Mikhail S. Burtsev (Dec. 8, 2022). *Recurrent Memory Transformer*. arXiv: 2207.06881 [cs]. URL: <http://arxiv.org/abs/2207.06881> (visited on 03/25/2025). Pre-published.
- Caron, Mathilde, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski, and Armand Joulin (May 24, 2021). *Emerging Properties in Self-Supervised Vision Transformers*. arXiv: 2104.14294 [cs]. URL: <http://arxiv.org/abs/2104.14294> (visited on 03/25/2025). Pre-published.
- Darcet, Timothée, Maxime Oquab, Julien Mairal, and Piotr Bojanowski (Apr. 12, 2024). *Vision Transformers Need Registers*. arXiv: 2309.16588 [cs]. URL: <http://arxiv.org/abs/2309.16588> (visited on 03/25/2025). Pre-published.
- Ilharco, Gabriel, Mitchell Wortsman, Ross Wightman, Cade Gordon, Nicholas Carlini, Rohan Taori, Achal Dave, Vaishaal Shankar, Hongseok Namkoong, John Miller, Hannaneh Hajishirzi, Ali Farhadi, and Ludwig Schmidt (July 2021). *OpenCLIP*. Version 0.1. Zenodo.
- Locatello, Francesco, Dirk Weissenborn, Thomas Unterthiner, Aravindh Mahendran, Georg Heigold, Jakob Uszkoreit, Alexey Dosovitskiy, and Thomas Kipf (2020). "Object-Centric Learning with Slot Attention". In: *Advances in Neural Information Processing Systems*. Vol. 33. Curran Associates, Inc., pp. 11525–11538.
- Lowe, David G. (Nov. 2004). "Distinctive Image Features from Scale-Invariant Keypoints". In: *IJCV* 60.2, pp. 91–110.
- Oquab, Maxime, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, Mahmoud Assran, Nicolas Ballas, Wojciech Galuba, Russell Howes, Po-Yao Huang, Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu Sharma, Gabriel Synnaeve, Hu Xu, Hervé Jégou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski (Feb. 2, 2024). *DINOv2: Learning Robust Visual Features without Supervision*. arXiv: 2304.07193 [cs]. URL: <http://arxiv.org/abs/2304.07193> (visited on 03/25/2025). Pre-published.
- Siméoni, Oriane, Gilles Puy, Huy V. Vo, Simon Roburin, Spyros Gidaris, Andrei Bursuc, Patrick Pérez, Renaud Marlet, and Jean Ponce (Sept. 29, 2021). *Localizing Objects with Self-Supervised Transformers and No Labels*. arXiv: 2109.14279 [cs]. URL: <http://arxiv.org/abs/2109.14279> (visited on 03/25/2025). Pre-published.
- Touvron, Hugo, Matthieu Cord, and Hervé Jégou (Apr. 14, 2022). *DeiT III: Revenge of the ViT*. arXiv: 2204.07118 [cs]. URL: <http://arxiv.org/abs/2204.07118> (visited on 03/25/2025). Pre-published.